

PROPOSAL: Acoustic UAS Classification

Safoan Miah

12 March 2026

1. Problem

The increase in use of unmanned aerial (drone) systems has created many airspace security problems, as seen in the 2018 Gatwick Airport shutdown and in the widespread commercial quadcopter use in the ongoing world conflict. Passive acoustic sensing is used alongside radar and radio frequency detectors since passive acoustic sensors are inexpensive, operate covertly, and do not need a direct line of sight or radio frequency emission from the target. However, real world use has three requirements that accuracy alone does not capture: generalization to new and unknown recording conditions, discrimination between drones and legitimate rotorcraft such as medical helicopters, tolerance of low signal-to-noise ratio (SNR).

This project will be a classification problem, classifying: drone, helicopter, background, where the helicopter class is included to test whether models can distinguish differences between acoustically similar legitimate rotorcraft rather than just dismissing them both into a single category.

2. Dataset

This project uses three public data sets. The Svanström Multi-Sensor Drone Detection dataset (Svanström et al., 2021) is the primary dataset and contains drone, helicopter and background recordings from three different Swedish airports. The Al-Emadi DroneAudioDataset (Al-Emadi et al., 2019) which includes additional drone audio to support the drone class and provides alternative sources for the cross-source generalisation evaluation. The ESC-50 (Piczak, 2015) dataset contains environmental noise to degrade the signal-to-noise ratio during testing; the Svanström set contains 90 ten-second recordings (30 per class) at 44.1 kHz; segmentation into 2-second sliding windows with 50% overlap giving ~810 labelled examples, resampled to 32 kHz mono to match the PANNs input specification. Class balance is somewhat uniform across drones, helicopters and backgrounds by design. Al-Emadi is reserved in full as the cross-source evaluation set and does not enter training. To avoid leakage between train/validation/test splits that could occur when using a small number of labeled audio examples, the data is split at the recording level, the audio is converted to 32 kHz mono to match the PANNs input specification, and then cached as log-mel spectrograms on Google Drive to keep per-epoch Colab memory and compute low.

3. Approaches

We will compare three architectures, which span across classical classification, convolutional learning to transfer learning. Drones produce a unique harmonic signature from the propellers frequencies which differs from helicopter and other environmental signatures (Al-Emadi et al., 2019); each model captures this differently. The first architecture makes use of a linear Support Vector Machine (SVM) over Mel Frequency Cepstral Coefficients (MFCC), Spectral Centroid, Roll Off, Band Width, Zero Crossing Rate and Root Mean Square Energy, providing a baseline performance based upon hand engineered features. Architecture 2 is a five-block CNN trained from scratch on

128-band log-mel spectrograms, following the spectrogram-as-image paradigm of Hershey et al. (2017); Casabianca and Zhang (2021) reported 94.7% mean accuracy on a comparable drone task using this approach. Architecture 3 uses PANNs CNN10 pretrained on AudioSet (Kong et al., 2020). PANNs are preferred over ImageNet-pretrained vision networks because they match the audio domain, placing this task in the small-data, similar-domain quadrant where fine-tuning is the indicated strategy: fine-tuned PANNs achieved 0.947 accuracy on ESC-50 while the same network trained from scratch achieved 0.833 (Kong et al., 2020).

Fine-tuning proceeds in three stages of progressive unfreezing, only the head is trained at learning rate 1×10^{-3} , the final convolutional block is unfrozen with differential learning rates (1×10^{-5} for convolutional layers, 1×10^{-3} for the head), finally the final two blocks are unfrozen on the same schedule. Macro F1 and per class AUC ROC are the primary metrics for measuring performance and are calculated with bootstrap 95% confidence interval. The three evaluation criteria are in distribution testing of the Svanstrom test set, cross source evaluation on the Al-Emadi hold-out (drone and background only), and an SNR degradation curve at clean (0dB), +10dB, +5dB, 0dB, -5dB using ESC-50 as the additive noise source.

4. Challenges

The limited number of samples is mitigated through the use of SpecAugment time and frequency masks (Park et al., 2019) and MixUp augmentation, both of which follow the same pipeline that was used to train PANNs (Kong et al., 2020). Leakage between windowed segments is prevented through the recording level split. The potential loss in performance due to a cross-source domain shift shown by Al-Emadi et al. (2019) to significantly decrease performance on unseen drone types is evaluated as part of the research question and not hidden. Over-fitting to the relatively small training set is handled through Dropout (0.3 to 0.5), L2 Weight Decay (1×10^{-4}), and Early Stopping based on Macro F1 Validation Performance. An alternative from fine tuning PANNs to feature extraction from frozen features is also available, where it achieved an accuracy of .918 on ESC-50 (Kong et al., 2020).

5. Timeline

Four weeks of active development will take place before the project deadline of April 24th, 2026. In week one (March 30-April 5) the datasets will be downloaded and preprocessed along with creating the recording level split manifests. In week two (April 6-12) the classical baseline model and CNN model will be trained and validated. In week three (April 13-19) three-stage PANNs fine-tuning, cross-source holdouts, and SNR curves with bootstrap intervals. And finally in week four (April 20-24) the draft will be completed, figures created, and checks for reproducing results with time to debug.

6. References

- Al-Emadi, S., Al-Ali, A., Mohammad, A., & Al-Ali, A. (2019). Audio Based Drone Detection and Identification using Deep Learning. *2019 15th International Wireless Communications & Mobile Computing Conference (IWCMC)*, 459–464.
<https://doi.org/10.1109/IWCMC.2019.8766732>
- Casabianca, P., & Zhang, Y. (2021). Acoustic-Based UAV Detection Using Late Fusion of Deep Neural Networks. *Drones*, 5(3), 54. <https://doi.org/10.3390/drones5030054>
- Hershey, S., Chaudhuri, S., Ellis, D. P. W., Gemmeke, J. F., Jansen, A., Moore, R. C., Plakal, M., Platt, D., Saurous, R. A., Seybold, B., Slaney, M., Weiss, R. J., & Wilson, K. (2017). *CNN Architectures for Large-Scale Audio Classification* (arXiv:1609.09430). arXiv.
<https://doi.org/10.48550/arXiv.1609.09430>
- Kong, Q., Cao, Y., Iqbal, T., Wang, Y., Wang, W., & Plumbley, M. D. (2020). *PANNS: Large-Scale Pretrained Audio Neural Networks for Audio Pattern Recognition* (arXiv:1912.10211). arXiv. <https://doi.org/10.48550/arXiv.1912.10211>
- Park, D. S., Chan, W., Zhang, Y., Chiu, C.-C., Zoph, B., Cubuk, E. D., & Le, Q. V. (2019). SpecAugment: A Simple Data Augmentation Method for Automatic Speech Recognition. *Interspeech 2019*, 2613–2617. <https://doi.org/10.21437/Interspeech.2019-2680>
- Piczak, K. J. (2015). ESC: Dataset for Environmental Sound Classification. *Proceedings of the 23rd ACM International Conference on Multimedia*, 1015–1018.
<https://doi.org/10.1145/2733373.2806390>
- Svanström, F., Alonso-Fernandez, F., & Englund, C. (2021). A dataset for multi-sensor drone detection. *Data in Brief*, 39, 107521. <https://doi.org/10.1016/j.dib.2021.107521>